

# CS2102 Notes.

## Introduction

Date

No.

### ACID Principles.

Atomicity, Consistency (no erroneous state)

Isolation (as if executions are sequential), Durability (lasting effects of executions).

### Equivalent Execution

Two executions are equivalent if they have the same effect on data.

### 3-Tier Architecture of DBMS.

- External, logical, physical

### Domain

A set of atomic values that an attribute can be, e.g. INT, TEXT. For a value of attribute  $A_i$ ,  $v \in \text{dom}(A_i)$  or  $v = \text{NULL}$ .

### Relation

A set of tuples. A relation is an instance of a schema, that defines a list of attributes and their domains, usually denoted as  $R(A_1, \dots, A_n)$ .

### Relation Database

A collection of tables (relations)

Integrity Constraints { Structural: key, foreign key, domain  
General: check, triggers. }

Superkey: subset of attributes that uniquely identifies a tuple.

Candidate key: a minimal superkey.

Primary key cannot be NULL

Foreign key: either contains **NULL** or appears as the primary key in the referenced relation

RA.

selection ( $\delta$ ):  $\delta[\text{condition}] R$

projection ( $\pi$ ):  $\pi[\text{col}_1, \text{col}_2] R$ : no duplicate cols, order matters, remove duplicate rows after  $\pi$ .

renaming ( $\rho$ ):  
①  $\rho[B_i \leftarrow A_i] R$  just renaming.  
②  $\rho[B_1, B_2, \dots] R$  reorder attributes.  
 $\Rightarrow$  need to write down all attributes.

Three-Valued Logic.

AND: True  $\wedge$  NULL  $\wedge$  False  $\rightarrow$  No False, but a NULL  $\Rightarrow$  NULL  
OR: True  $\vee$  NULL  $\vee$  False.  $\rightarrow$  No False or NULL  $\Rightarrow$  True

(NULL  $>$ ,  $<$ ,  $\leq$ ,  $\geq$ ,  $=$ ,  $\neq$ ,  $+ - \times \div A$ )  $\Rightarrow$  NULL.

However,  $\text{NULL} \equiv \text{NULL} \Rightarrow \text{True}$ ,  $\text{NULL} \equiv A \Rightarrow \text{False}$ .

$$A \equiv B \Rightarrow A = B$$

'Union Compatible': R and S have the same number of attributes and each corresponding attribute has the same or compatible domains.

Cross Product:  $R \times S$ : R and S, unlike SQL, must be disjoint.

JOIN:

Inner  $\left\{ \begin{array}{l} \theta\text{-Join: } R \bowtie_{[\theta]} S = \delta_{[\theta]}(R \times S) \\ \text{Equi-Join: } \theta \text{ is based on equality } (= \text{ or } \equiv) \\ \text{Natural Join: } R \bowtie S \Rightarrow R[A_1, \dots, A_i, \underbrace{B_1, \dots, B_j, C_1, \dots, C_k}_{\substack{\text{order sometimes matters} \\ \text{only } \in R}}] \times S[C_1, \dots, C_k, \underbrace{D_1, \dots, D_l}_{\substack{\text{order sometimes matters} \\ \text{only } \in S}}] \end{array} \right.$

Outer { Left Join:  $D[0]$  include dangling tuples from left-right  
 Right Join  $D[0]$  ...  
 Full Join  $D[0]$  ...  
 } Pad missing values with NULL  
 May or may not be natural.

### Strong VS Weak Equivalence.

- A SE B  $\Leftrightarrow$  (A and B produce error) or the same result
- A WE B  $\Leftrightarrow$  (A or B produces error) or the same result

### Principle of Acceptance.

- Perform only when the condition is True.
- Used in WHERE clause and RA . and PL/SQL IF statement

### Principle of Rejection.

- Reject only when the condition is False . (perform otherwise)
- Used in integrity constraints .

## SQL.

### DDL Data Definition

CREATE, ALTER, DROP, ADD

### DML Data Manipulation

INSERT, UPDATE, DELETE

### DQL Data Query

SELECT

### DCL Data Control (not covered).

### TCL Transaction Control

BEGIN, COMMIT, ROLLBACK.

SQL is case-insensitive / INSERT operation is atomic if there are multiple rows . Missing values are replaced with NULL or default value . / Principle of Acceptance is applied to DELETE ( delete if true ) . / Use CONSTRAINT name to name a constraint .

Use Principle of Rejection to check UNIQUE constraint  
 $\text{UNIQUE} \Leftrightarrow \text{expr} \rightarrow \text{true or } \text{NULL}$

Foreign Key Constraint. → Set on the referencing table.  
 NO ACTION, RESTRICT (= NO ACTION + non-deferred check), CASCADE, SET DEFAULT (the default value must be there in the referenced table), SET NULL.

INITIALLY IMMEDIATE is always IMMEDIATE if the transaction does not set it to DEFERRED. Even if it does, the effect<sup>only</sup> lasts within the scope of it.

Query. → at most one NULL returned.

`SELECT [DISTINCT] <attr-list> FROM <rel-list>  
 WHERE <conditions>`

Operators: / : sqrt, || concat.

Pattern Matching: - : any char %: string of 0 or more char

SET operations: ALL keeps duplicates: every value has a count. Without ALL, normal set operations are done. UNION/INTERSECT/EXCEPT.

JOIN without NATURAL, whether INNER or OUTER, should be accompanied by ON [conditions].

Universal Schema Assumption.

All tables can be represented by a (possibly very wide) table. → May not always hold.  
 selected cannot be NULL

can be replaced by  
 inner join  
 outer join

IN: ( $\exists \text{row} \in \text{subquery} : \text{row} = \langle \text{expr} \rangle$ ) evaluates to True.  
 NOT IN: ( $\text{row} : \text{row} \neq \langle \text{expr} \rangle$ ) evaluates to True.  
 i. It is possible that both return nothing. (esp. when there is NULL in <subquery>)

EXISTS: <subquery> returns at least one row  
NOT EXISTS: otherwise.

In the jargon of querying, "known"  $\rightarrow$  NOT NULL.

Date

No.

ANY:  $\exists \text{ row} : \langle \text{expr} \rangle \text{ row} = \text{True}$ . [ $\langle \text{subquery} \rangle$  must have only one col.]  
ALL:  $\forall \text{ row} : \langle \text{expr} \rangle \text{ row} = \text{True}$ .

$\text{IN} \equiv (\subseteq \text{ANY}) / \langle \text{el} \rangle \text{ cop} \text{ ANY } \langle \text{Q} \rangle \equiv \text{EXISTS} (\dots \text{ AND } \langle \text{el} \rangle \text{ cop} \text{ } \langle \text{el} \rangle \text{ cop})$   
 $\text{NOT IN} \equiv (\subsetneq \text{ALL})$ .

If duplicate removal is needed for SELECT, then the attr. by which the table is sorted must be also in the SELECT clause

\* Smaller lexicographic order has higher rank  $\rightarrow$  Asc

## Aggregation Functions

- Only applicable to non-NUL values, except \*.
- Can have DISTINCT options.
- If all values are NULL, MAX(A)  $\Rightarrow$  NULL, COUNT(A)  $\Rightarrow$  0.
  - COUNT(\*)  $\Rightarrow$  n.

## Grouping

- t and t' will be in the same group if  $(\bigwedge_i t.a_i \text{ IS NOT DISTINCT FROM } t'.a_i)$
- Restrictions: If AER is in the SELECT clause then.
  1. A is in the GROUP BY clause, or
  2. A is the input for aggregation function in the SELECT clause, or
  3. The primary key of R is in the GROUP BY clause.

## HAVING

- Similar to WHERE, both using PofA.
- Restrictions: just the same as above except the blue part.

## Evaluation Order

From  $\rightarrow$  WHERE  $\rightarrow$  GROUP BY  $\rightarrow$  HAVING  $\rightarrow$  SELECT  $\rightarrow$  ORDER BY  $\rightarrow$  LIMIT/OFFSET.

CASE  
 WHEN <exp> THEN <val>  
 ELSE <default>  
 END.

CASE <exp>  
 WHEN <result> THEN <val>  
 ELSE <default>  
 END.



COALESCE → return the first non-NULL value.  
 $(\text{NULL}, (v = u)) \Rightarrow \text{True}$ .

NULLIF(v, u) =  $\begin{cases} v, & \text{otherwise} \\ \text{NULL}, & (v = u) \end{cases}$

Common Table Expressions (At most 2 in exams!)

WITH E(attr<sub>1</sub>, ..., attr<sub>n</sub>),  
 CTE-1 AS (Q-1)

View:

CREATE VIEW <name> AS (Q).

Universal Quantification.

• Double Negation Solution.

$\forall x : P(x) \Leftrightarrow \neg \neg \forall x : P(x) \Leftrightarrow \neg \exists x : \neg P(x)$ .

• Cardinality Solution.

Use set operations and COUNT(P DISTINCT  $\ast$ ).

Recursive Query.

• WITH RECURSIVE ... AS (

... count

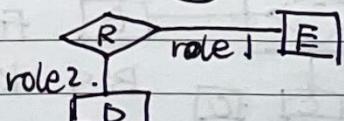
UNION [ALL] .SELECT ... count+1 FROM L.  
 ... ).

• Can have lazy evaluation: move the termination condition from CTE to the main query. (WHERE count < a)

Entity Sets  $\square$  cannot be proper nouns.

Attributes: (1) key  $\circlearrowleft$  attr (2) composite  
 (3) multivalued  $\circlearrowleft$  A (4) derived  $\circlearrowleft$  A

Relationship sets



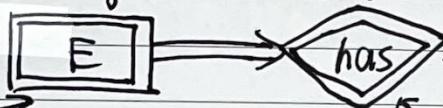
## Relationship Constraints.

- Cardinality Constraint: many/one -to - many/one.
  - Key Constraint: at most one relationship per entity.
- Participation Constraint:
  - Partial Participation: lower limit = 0.
  - Total Participation: lower limit = 1.

key + Total Participation  $\Rightarrow$

One and only one relationship instance per entity.

## Dependency Constraint.



D  $\leftarrow$  Owner Entity

Identifying Relationship sets.

Weak Entity Sets

Only have partial keys.

## Multivalued Attributes Mapping.

- ① Split to columns  $\Rightarrow$  limited number of possible values.
- ② A separate table  $\Rightarrow$  query 2 tables at a time.

## ER Mapping.

Many-to-Many : Set a table with 2 foreign keys.

Many-to-One : The "one" table references the "many" table

One-to-One : Both tables have a unique FK (but no "NOT NULL").

Key + Total : set the foreign key to be UNIQUE and NOT NULL

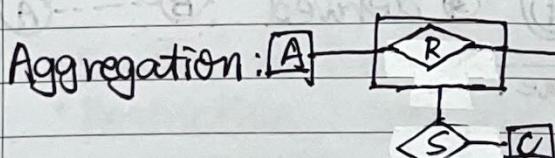
Weak Entities : plus the key of the owner.

## ISA Hierarchies.

Covering Constraint		
Overlap	False	True $\Rightarrow$ superclass belongs to at least one subclass
F	↓	↓
T		

superclass can belong to multiple subclasses.

- Remember to do ON DELETE CASCADE in subclasses.



for mapping, use A's and B's keys in S.

# SQL Programming

## Dynamic SQL:

- EXEC SQL
  - PREPARE var FROM :stmt;
  - EXECUTE var USING para1, para2; // ← parameters
  - DEALLOCATE PREPARE var;

Function

```

CREATE OR REPLACE FUNCTION func([IN/OUT] p1, [IN/OUT]
p2 ... ) one value a tuple a set of tuples
RETURNS TYPE/TABLE_NAME/SETOF TABLE/SETOF RECORD/
TABLE (attr1, attr2 ...)/TRIGGER AS $$.
DECLARE ... ...
BEGIN ... ...
END; $$ LANGUAGE plpgsql/c/sql/...
    
```

Use CALL for procedures and SELECT for functions.

Cursor

```

curs CURSOR FOR SELECT ...; OPEN curs; LOOP.
FETCH curs INTO r; EXIT WHEN NOT FOUND;
[RETURN NEXT] /* not curs' next */; END LOOP; CLOSE curs;
has something to do with the function signature.
    
```

\* FETCH [PRIOR/FIRST/LAST/ABSOLUTE n] FROM cur INTO r;

## SQL Injection Attacks

- Malicious instructions are inserted as input.
- Approach ①: Use function/procedure to isolate the input from the execution. → the input is regarded as a string constant.
- Approach ②: PREPARE can be OR

Trigger

```

CREATE TRIGGER name AFTER/BEFORE/INSTEAD OF
INSERT/UPDATE/DELETE ON tname FOR EACH ROW/
STATEMENT [WHEN (cond)] EXECUTE FUNCTION func()
[DEFERRABLE...].
    
```

Special variables a trigger can access:

- TG\_OP: the operation that triggers
- TG\_TABLE\_NAME
- OLD / NEW

## Return Values

	Has effect				Has no effect		
	INSERT	UPDATE	DELETE		INSERT	UPDATE	DELETE
BEFORE	non-NULL	non-NULL	non-NULL	New	OLD/NULL	OLD/NULL	NULL
AFTER	NO	NEW	OLD		Any		

## INSTEAD OF

- Instead of doing something on a VIEW, do it somewhere else, usually the corresponding table.
- E.g., instead of modifying the current mark-score VIEW, update the table directly.
- Always row-level. ; • Return NULL to skip the original op.

## Statement Level.

- No access to NEW, OLD.
- Ignores return values.
- Raise an exception to skip subsequent operations.
- Can be BEFORE.

## Deferred Trigger.

- Only works with AFTER and FOR EACH ROW
- Activated by COMMIT.

## Multiple Triggers.

- Order of activation
  - 1. BEFORE statement-level
  - 2. BEFORE row-level
  - 3. AFTER row-level
  - 4. AFTER statement-level
- If a BEFORE row-level returns NULL, then all subsequent triggers are skipped.

## Normal Forms

- Reduces data redundancy and improves data integrity.

## Redundancy $\rightarrow$ Anomalies.

- Update anomaly: partial update that does not cover all rows supposed to be updated.
- Insertion anomaly: inability to add data due to absence of other data, e.g., fk or pk.
- Deletion anomaly: cannot delete a data because it contains something that cannot be deleted, e.g., pk.

## Functional Dependencies.

- $A \rightarrow B \Leftrightarrow$  if A is the same, then B is the same.
- An FD may hold on table A but not on B.
- Huang's Theorem: Let M, N be the set of rows uniquely identifiable by  $A \setminus \{A, B\}$ , resp., then  $(A \rightarrow B) \Rightarrow M = N$ .

- Armstrong's Axioms.
- Reflexivity  $AB \rightarrow A$
  - Augmentation  $(A \rightarrow B) \Rightarrow (AC \rightarrow BC)$
  - Transitivity  $(A \rightarrow B \wedge B \rightarrow C) \Rightarrow (A \rightarrow C)$

## Additional Rules

Date \_\_\_\_\_

No. \_\_\_\_\_

Rule of Decomposition:  $(A \rightarrow BC) \Rightarrow (A \rightarrow B \wedge A \rightarrow C)$

Rule of Union:  $(A \rightarrow B \wedge A \rightarrow C) \Rightarrow (A \rightarrow BC)$

## Closure

- $\{A\}^+ = \{ABC\dots\}$  = the set of attributes that can be determined by A, directly or indirectly.
- Traverse every unused FD until no new attribute is added in a loop.

## Finding Keys.

- List out all possible combinations of attributes and their closures  $\rightarrow$  If a closure is complete  $\Rightarrow$  superkey
- Trick ①: attributes that do not appear in any RHS are definitely part of the key.
- Trick ②: try small attributes first. If a set of attributes is a superkey, so is its superset.
- Trick ③: if  $\{AB\}$  is a key, then other keys do not completely contain  $\{AB\}$   $\Rightarrow$  stop checking a set if it contains  $\{AB\}$ .

## Prime Attribute:

- An attribute is a prime attribute if it appears in a key.

## Soundness & Completeness.

- Sound  $\Leftrightarrow$  Everything provable is true.
- Complete  $\Leftrightarrow$  Everything true has a proof.

## Non-Trivial & Decomposed FD.

- Non-trivial: RHS does not appear in LHS.
- Decomposed: RHS only has one attribute.
- Find them using closures

## BCNF

- A table is in BCNF if every non-trivial and decomposed FD has a superkey as its LHS.

BCNF intuition:

- Every attribute should depend only on superkeys.
- Any dependency on non-superkeys is prohibited.

Tricks: For a table of  $n$  attributes, only need to check "more-but-not-all" closures of  $\leq n-2$  attributes

## OK to BCNF Algorithm

- Consider explicit FDs only
- ① Check the closure of each attribute subset.
  - ② Check if there is a "more-but-not-all" closure.
  - ③ Start from smaller sets. Skip sets that contain a confirmed superkey.

## BCNF Decomposition

- ① Input a table  $R$ . Suppose  $\{X\} \rightarrow \{X\}^+$  is a "more-but-not-all" closure.
- ② Split  $R$  into  $R_1 = \{X\}^+$  and  $R_2 = R - \{X\}^+$  → no repeated rows.
- ③ Further check  $R_1$  and  $R_2$ .
  - Trick: If there are multiple choice of  $\{X\}$ , pick the smallest  $\{X\}^+$ . If  $|R_1| = 2$ , then  $R_1$  must be in BCNF

## FD Projection on $R_i$

- can have multiple attributes.
- ① Enumerate all attribute subsets on  $R_i$
  - ② Derive the closures on  $R_i$ .
  - ③ Remove attributes in the closure but not on  $R_i$ .
    - If there is still a "more-but-not-all" closure, then keep decomposing. → using closure!!
    - If the union of all projected FDs is equivalent to the original set of FDs, then we say that the decomposition is FD-preserving.

## Properties of BCNF:

Good: No anomalies (update/insert/delete), small redundancy, reconstructable

Bad: Not FD-preserving.

Lossless Join:  $R_1 \bowtie R_2 = R$ . Lossy Join  $R \subsetneq R_1 \bowtie R_2$   
A decomposition is lossless iff the full set of common attributes in  $R_1$  and  $R_2$  constitute a superkey of  $R_1$  or  $R_2$ .

→ Two sets of FDs  $S, S'$  are equivalent iff:

Every  $FD \in S$  can be derived from  $S'$  and the other way around. → Using closure is efficient.  $S$  and  $S'$  should have the same closure.

Properties of 3NF: Lossless join, FD-preserving, small redundancy.  
3NF: Every non-trivial and decomposed FD: LHS is a superkey or RHS is a prime attribute.

BCNF is stricter  $\Rightarrow$   $\text{BCNF} \Rightarrow 3NF$  but  $3NF \neq BCNF$ .

3NF check: For each given FD, check LHS & RHS.

3NF Decomposition Algorithm / 3NF Synthesis

- ① Derive a minimal basis
- ② Combine FDs whose LHSs are the same.
- ③ Create a table for each FD remained.
- ④ If none of the tables contains a key, then create a table that contains a key  $\rightarrow$  Ensures lossless join
- ⑤ Remove redundant tables.

Minimal Basis/Cover, M, for a set of FDs, S.

- M and S are equivalent.
- M contains only non-trivial and decomposed FD.
- If any FD is removed from M, then M cannot derive S.
- For every FD  $A \rightarrow M$ , if we remove an attribute from LHS, then it cannot be derived from S.

Algorithm to get Minimal Basis

- ① Transform each FD such that each RHS contains only one attribute.  $\rightarrow$  the current working set.
- ② Remove redundant attribute from each FD. E.g.  $AB \rightarrow C$ . If  $A \rightarrow C$  is implied by M, then remove B.
- ③ Remove redundant FD if it can be derived from the rest of M.  $\rightarrow$  Again, use closures.

Appendix.

Non-Trivial FDs:

$A \rightarrow B$  if  $B \not\subseteq A$ . This means  $\{A, B\} \rightarrow \{A\}$  is trivial

Completely non-trivial if  $A \cap B = \emptyset$

Armstrong's Axioms Tricks.

- $A \rightarrow B, A \rightarrow C ; A \rightarrow AB \rightarrow A \rightarrow AC \Rightarrow ABB \rightarrow ABC \rightarrow BC$   
 $\therefore A \rightarrow BC$ .

- Augment given FDs to some steps obtained.